

Benchmarking the state-of-the-art Task Execution Frameworks of Many- Task Computing

Thomas Dubucq - Tony Forlini - Virgile Landeiro Dos Reis - Isabelle Santos

Introduction

▶ Systems:

- ▶ Charm++
- ▶ HPX
- ▶ Legion
- ▶ MATRIX
- ▶ Sparrow
- ▶ Swift/T

▶ Testbeds:

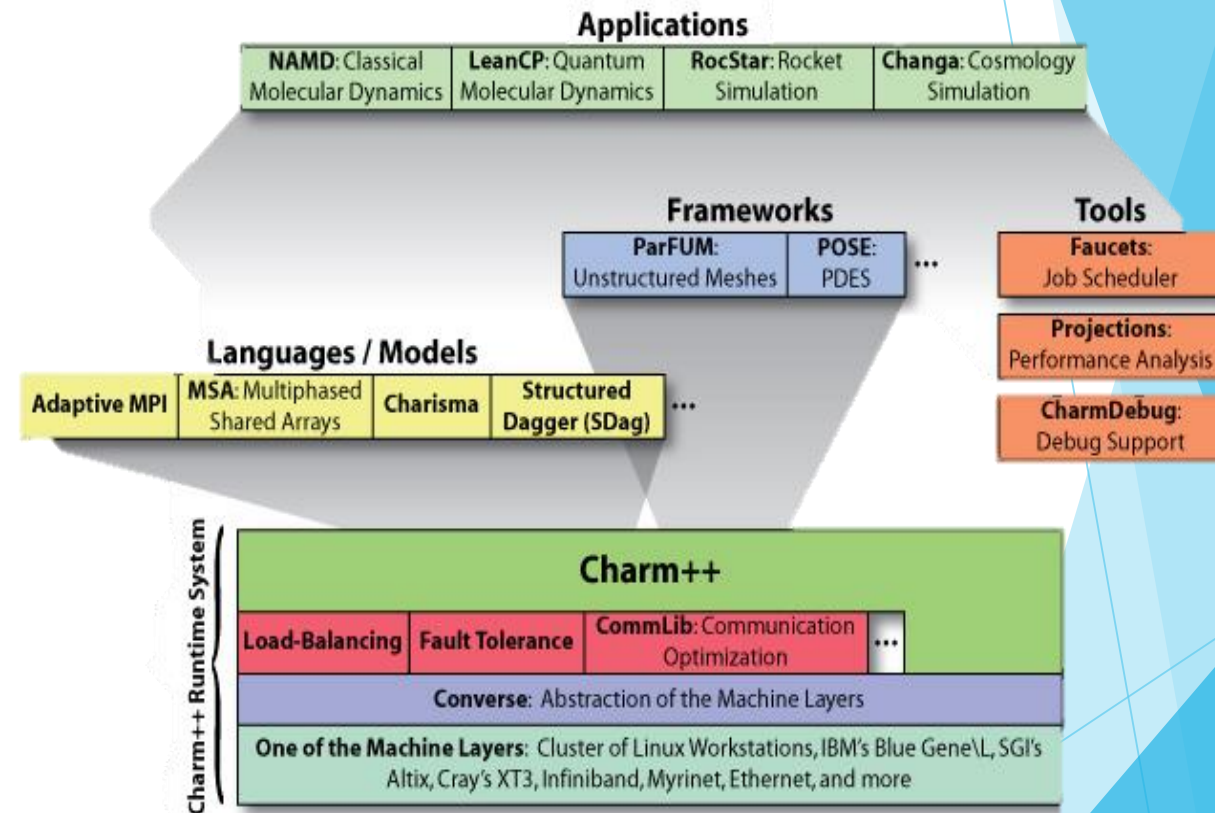
- ▶ From 1 to 128 nodes
- ▶ C3.large instances
- ▶ 2 cores per node
- ▶ Weak scaling

▶ Metrics:

- ▶ Throughput
- ▶ Efficiency
- ▶ Task latency

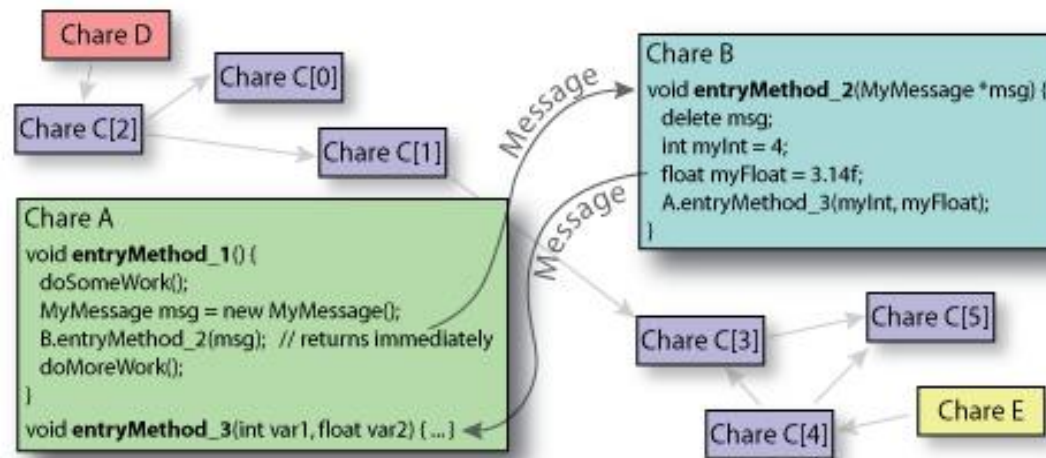
Charm++

- ▶ Explicit parallelism
- ▶ Applications:
 - ▶ NAMD molecular dynamics
 - ▶ LeanCP quantum molecular dynamics
- ▶ Tools:
 - ▶ Faucets job scheduler
 - ▶ CharmDebug
 - ▶ Projections

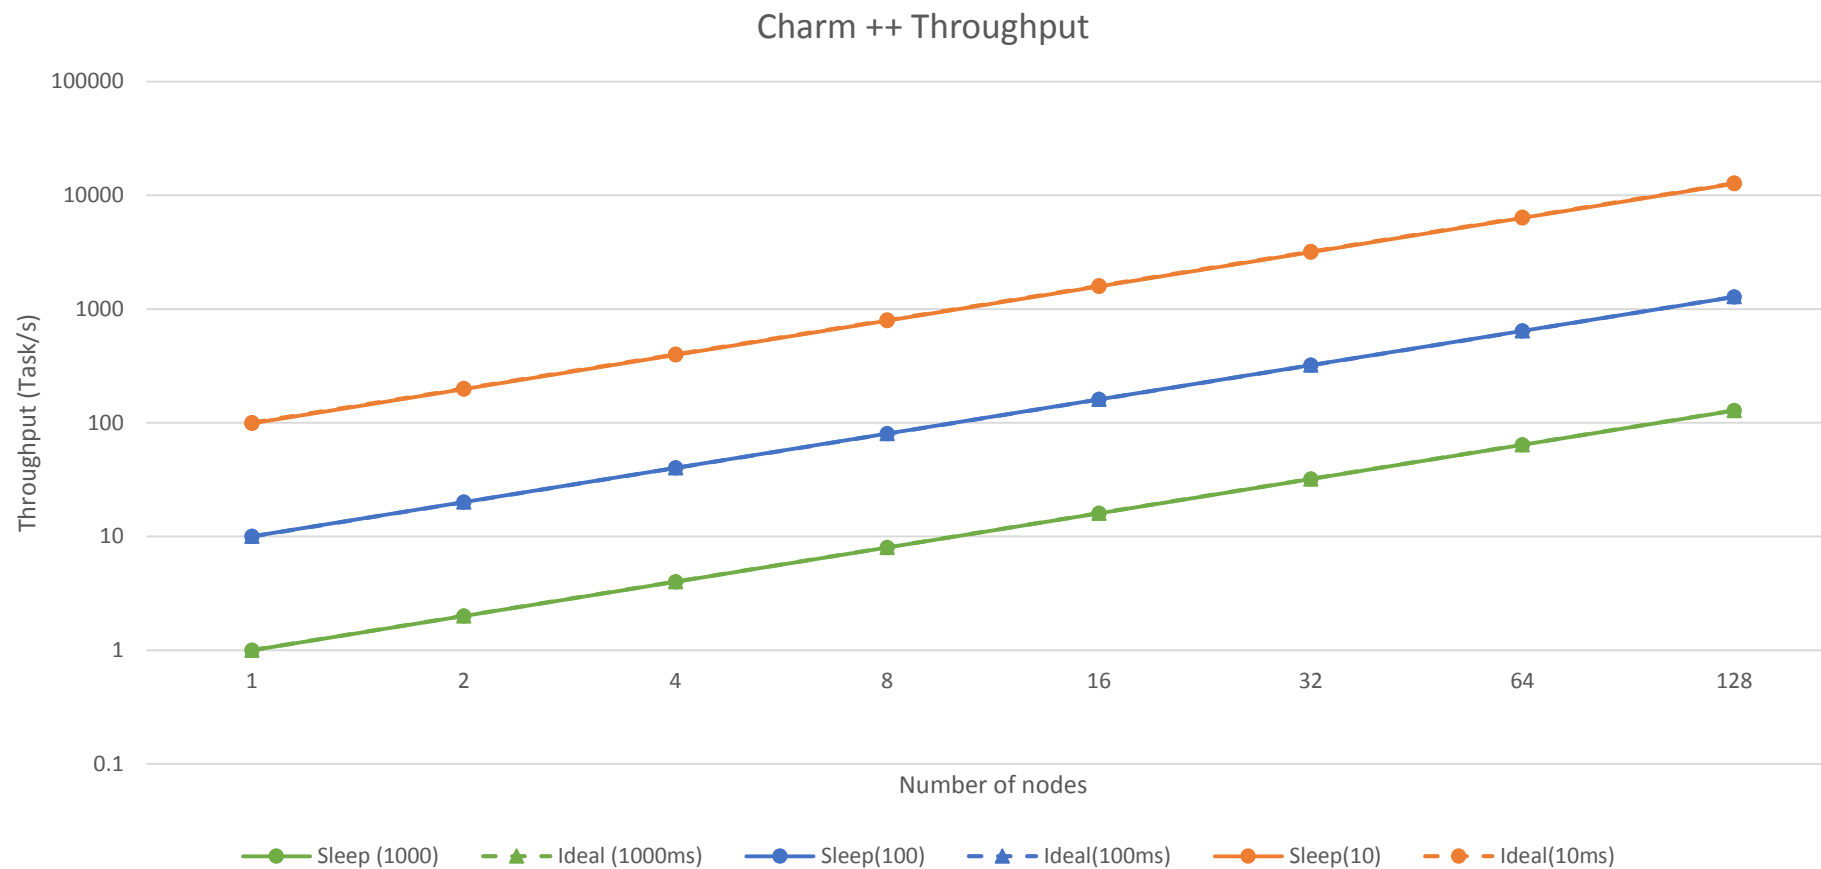


Charm++


- ▶ Asynchronous message passing parallel programming paradigm
- ▶ Tasks defined as chare objects
- ▶ Chare objects communicate with each other by sending messages




Charm++ - Results



Charm++ - Efficiency

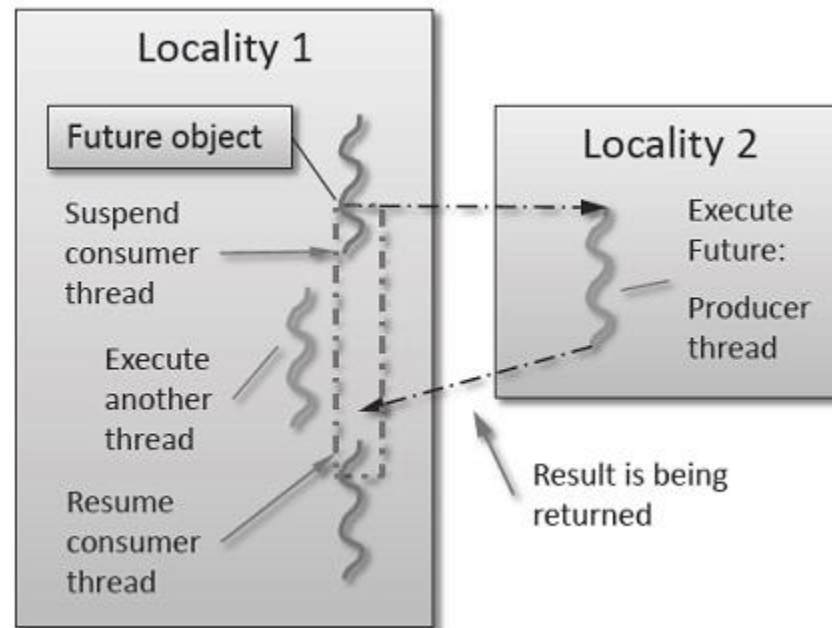


Charm++ - Results



HPX (High Performance ParallelX)

- ▶ Boost C++ runtime system for parallel and distributed applications
- ▶ ParalleX execution model
 - ▶ hide latency by switching tasks
 - ▶ fine-grained tasks
 - ▶ reduce overhead
- ▶ Futures
 - ▶ delayed computation
 - ▶ reader thread & producer thread



HPX: Challenges

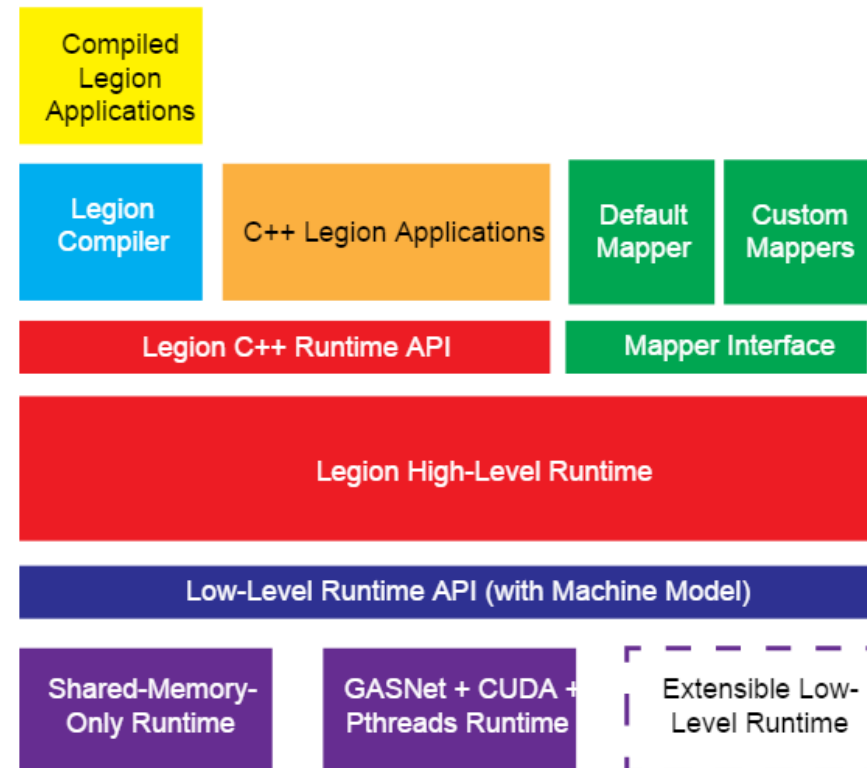
- ▶ Never previously deployed on EC2
- ▶ Incomplete documentation
 - ▶ unmentioned dependencies
 - ▶ typos in command lines
- ▶ Long compilation time
- ▶ Memory requirements
- ▶ Version compatibility among dependencies

Legion


- ▶ Programming model for heterogeneous, distributed machines
- ▶ Built on top of GASNet communication system
- ▶ Heterogeneous
 - ▶ Mixed CPUs and GPUs
- ▶ Distributed
 - ▶ Large spread and variability of communication latencies
- ▶ Implicit parallelism

Legion


System Architecture



Legion - Results

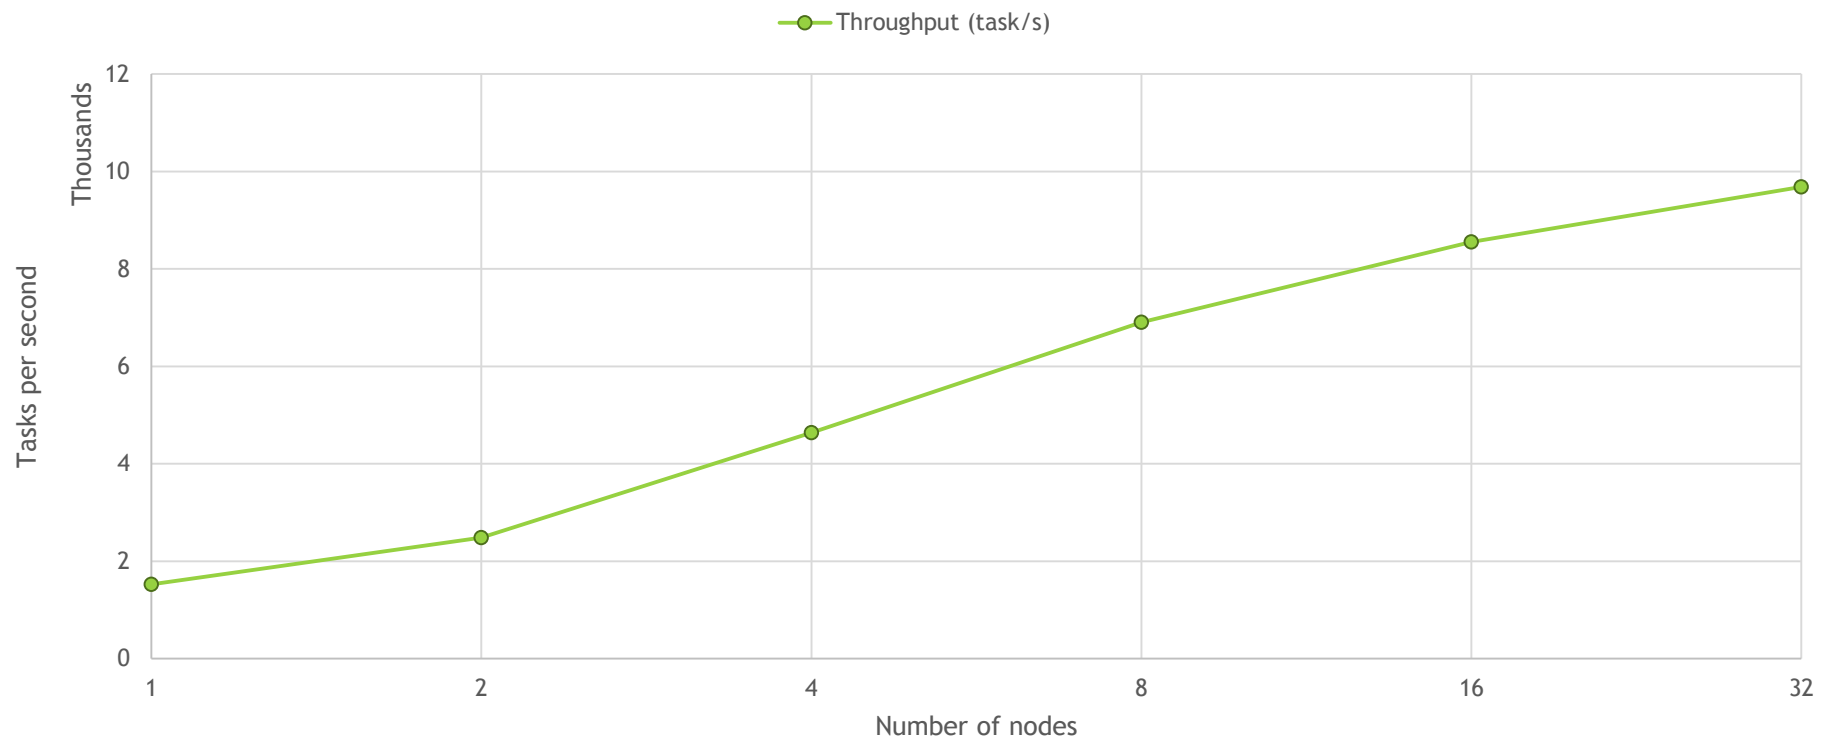


Legion - Results



Legion - Results


Throughput (tasks/s) for sleep(0) Tasks



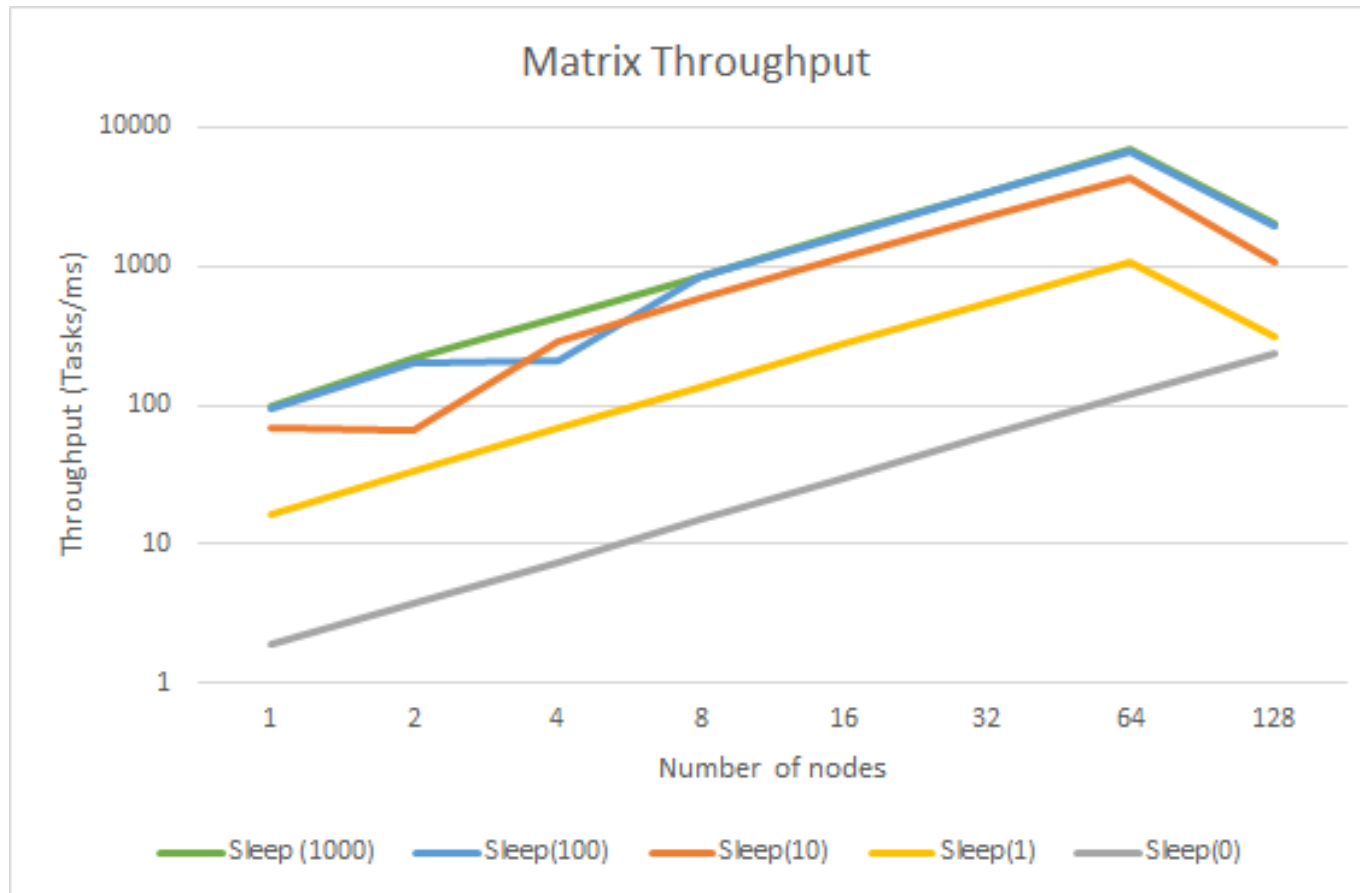
MATRIX

- ▶ Distributed data-aware execution fabric
- ▶ Work stealing
- ▶ Built on top of ZHT
- ▶ Supports MTC and HPC workloads
- ▶ Coded in C++

MATRIX - Results

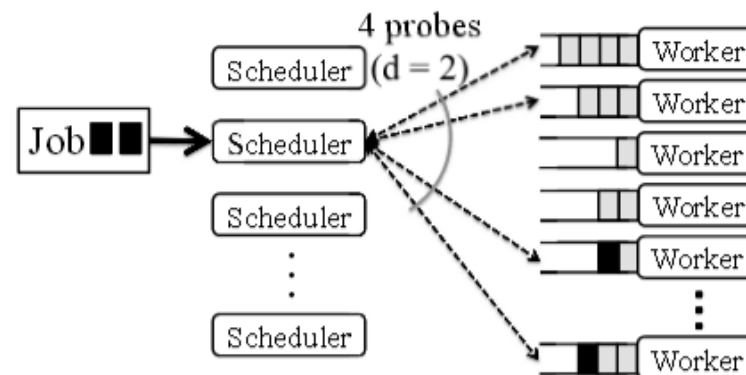
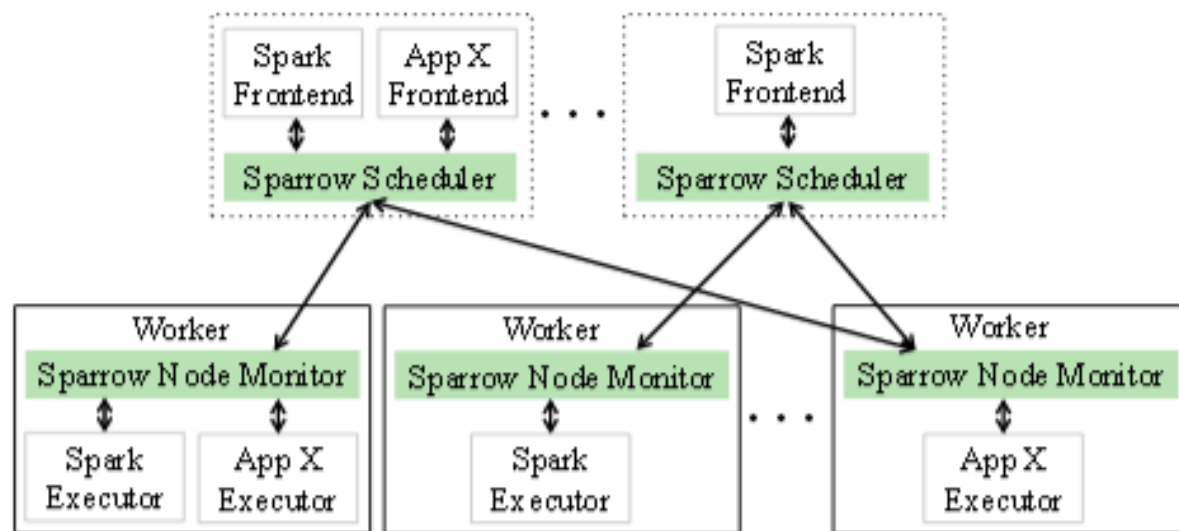


MATRIX - Results



Sparrow

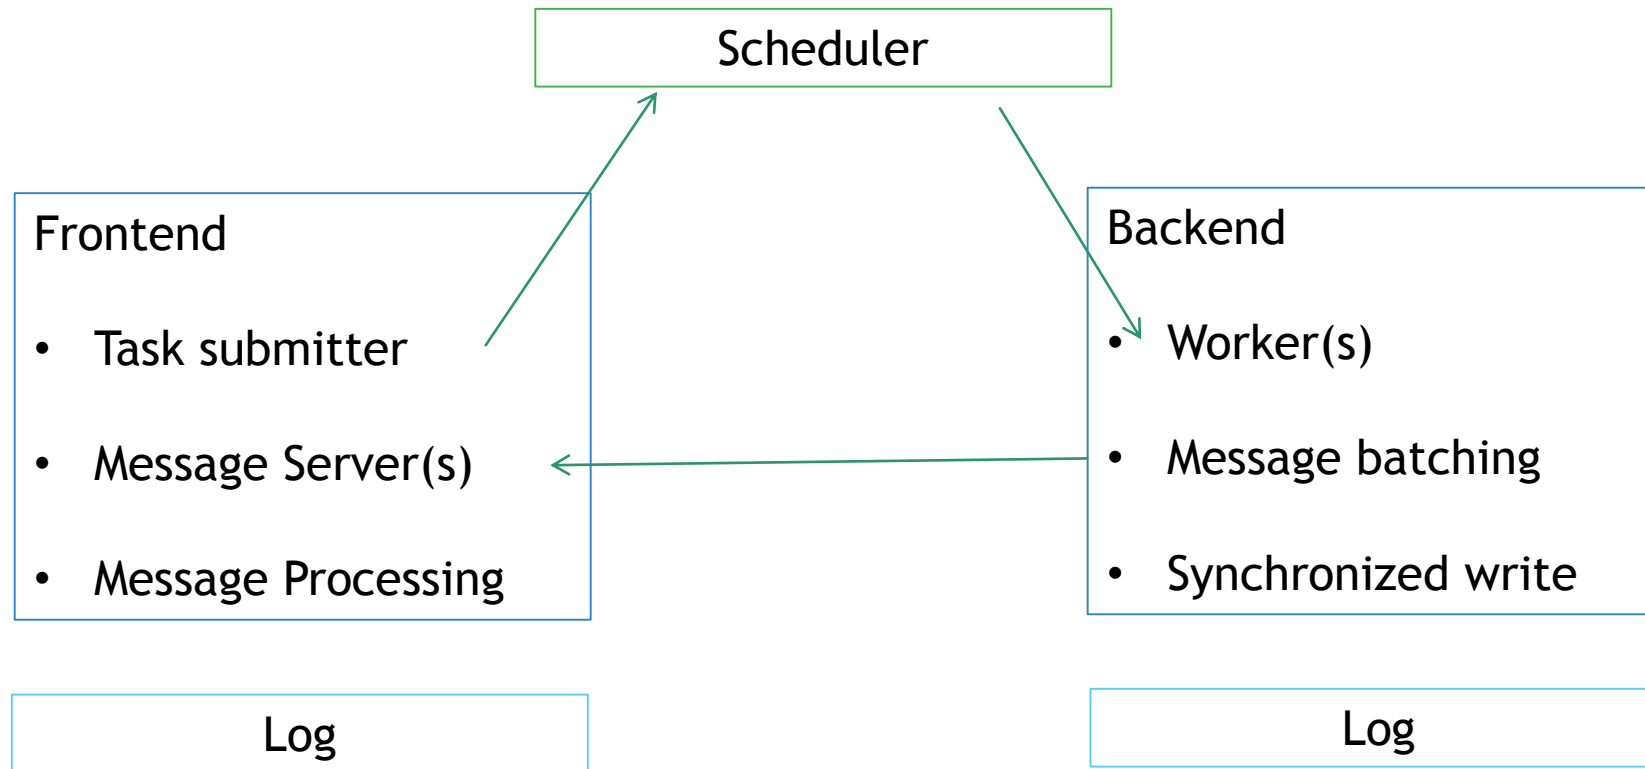
- ▶ Stateless distributed scheduler
- ▶ Late binding
- ▶ Batch sampling
- ▶ Coded in Java




Sparrow - Challenges

- ▶ Backend and Frontend provided in the code source don't match the benchmark metrics
- ▶ Necessity of task completion acknowledgement
- ▶ Dense workload handling


Sparrow



Sparrow - Results

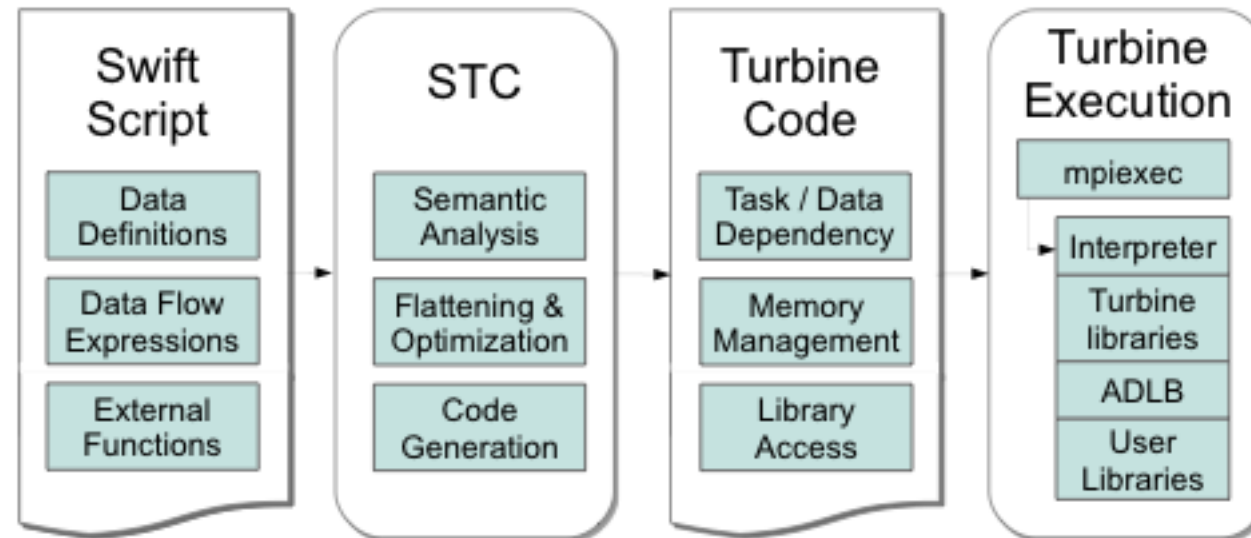


Sparrow - Results



Swift/T

- ▶ Workflow system
- ▶ Built on top of MPI
- ▶ Highly programmable




- ▶ Improve from Swift/K by introducing a parallel evaluation of complex scripts
- ▶ Two main tools:
 - ▶ STC: Java-based compiler that produces Turbine code
 - ▶ Turbine: execution combining MPI and ADLB

Swift/T - Challenges

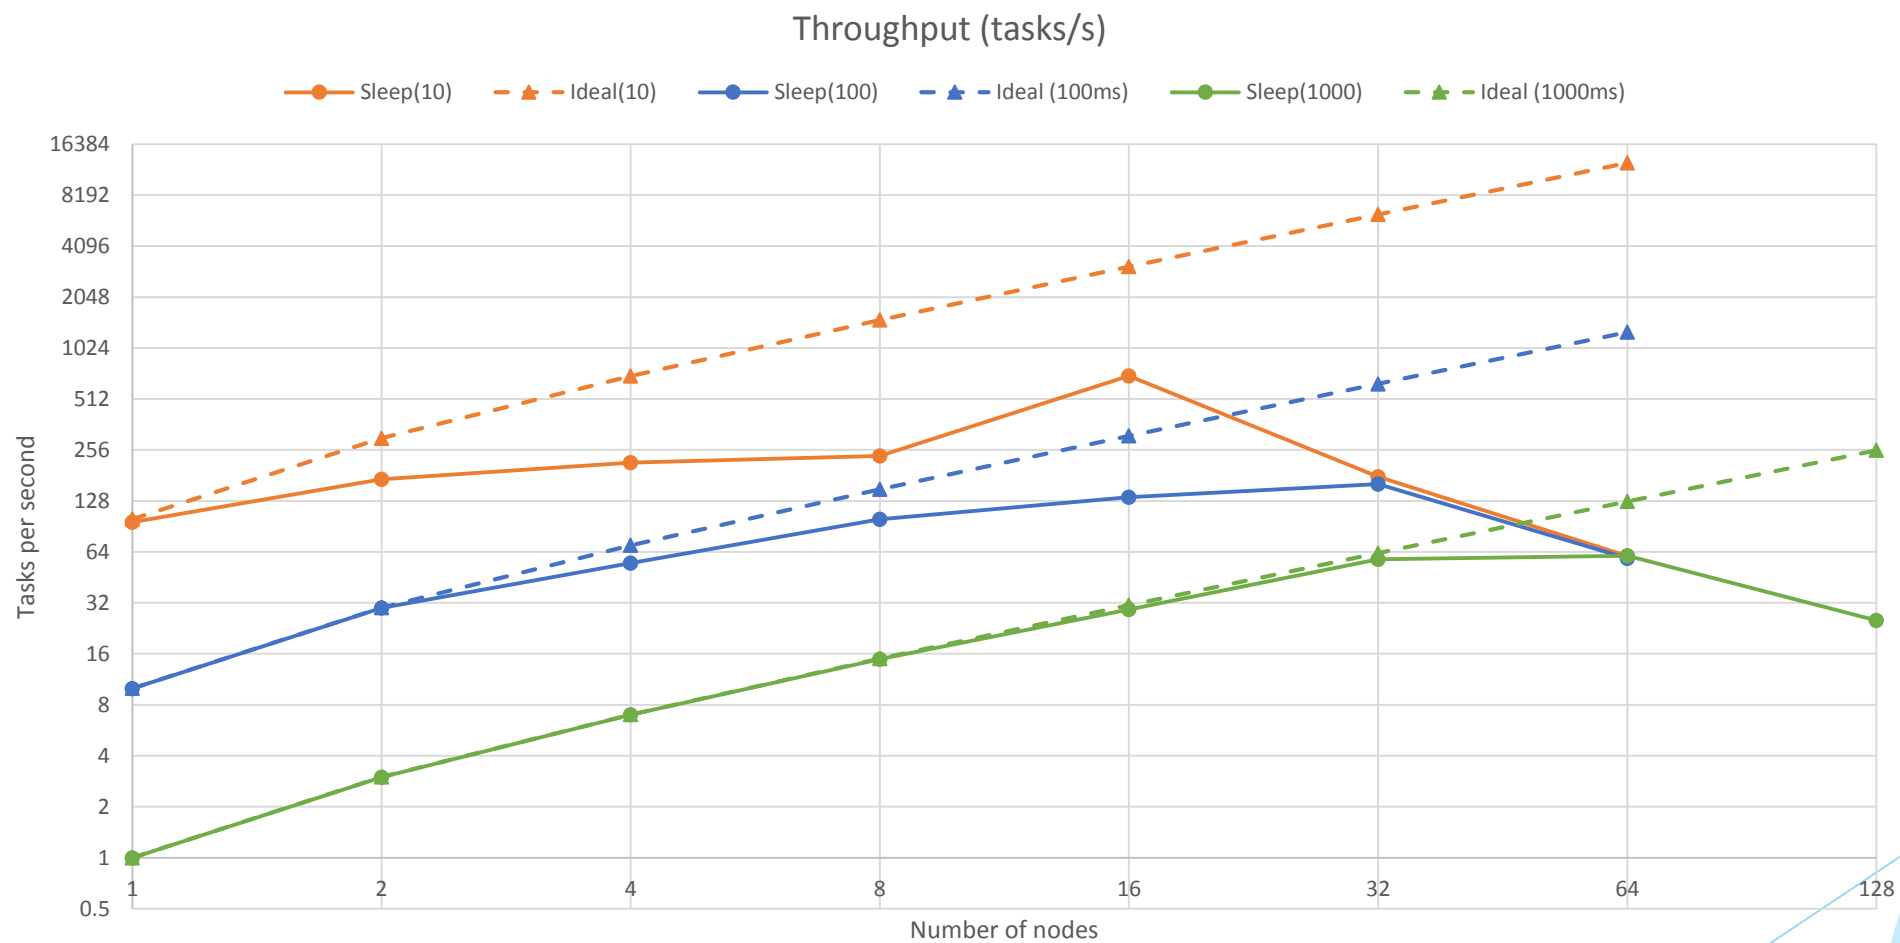
- ▶ Very sparse documentation to deploy on EC2 but very reactive support by Justin M. Wozniak.
- ▶ Sleep(0) tasks do not scale. Have not found why yet but suspect the following piece of code:

```
@dispatch=WORKER
(void v) sleep(float seconds) "turbine" "0.0.4" "sleep" [
    "if { <<seconds>> > 0 } { after [ expr {round(<<seconds>> * 1000)} ] }"
];
```


Throughput for sleep(0) (tasks/s)



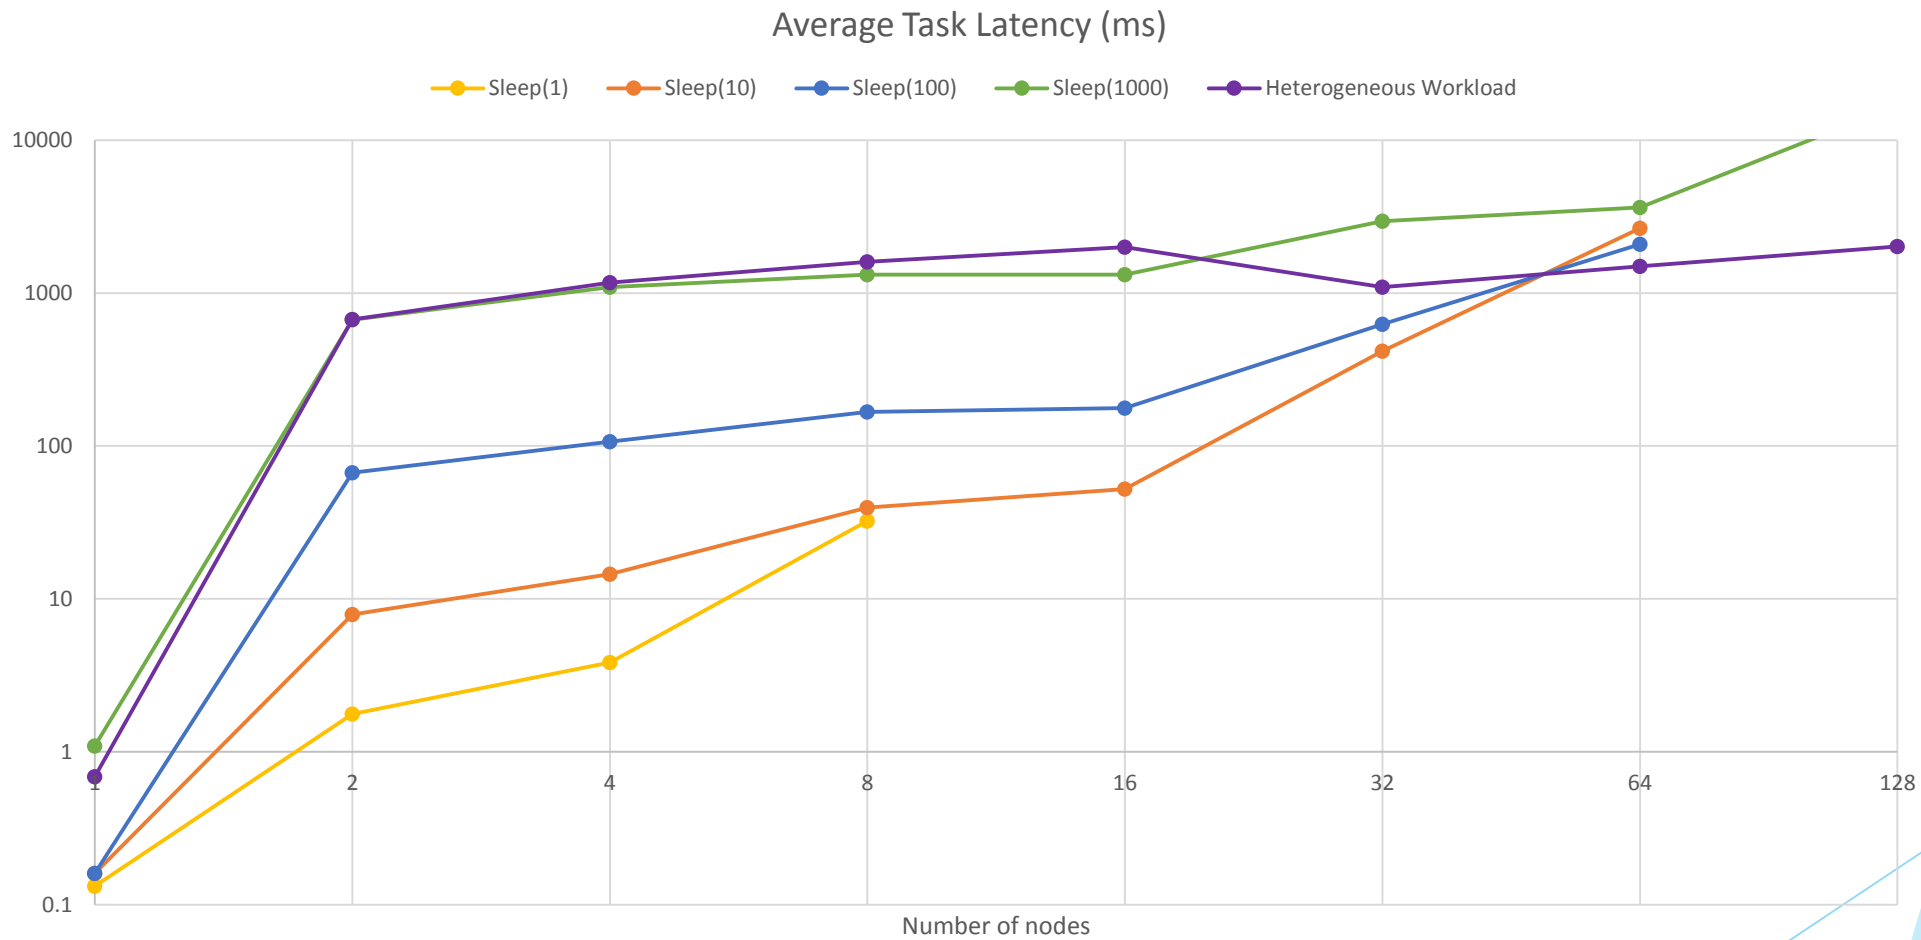
Swift/T - Results



Swift/T - Results




Swift/T - Results



Conclusion

- ▶ Performance ranking:

- ▶ Charm++
 - ▶ Legion
 - ▶ Sparrow
 - ▶ MATRIX
 - ▶ Swift/T
- 
- best
- worst

- ▶ Non evaluated systems:

- ▶ HPX